

```
1
2 // PLEASE COMPILE WITH g++ reduce_par.cpp -lpthread
3
4 #include <memory>
5 #include <iostream>
6 #include <vector>
7 #include <thread>
8 #include <cassert>
9 #include <string>
10 // #include <numeric>
11 // #include <execution>
12 #include <parallel/numeric>
13
14 template <typename T>
15 T sum(T a, T b) {
16     return a + b;
17 }
18
19 template <typename T>
20 T max(T a, T b) {
21     return a > b ? a : b;
22 }
23
24 // template <typename T, typename = std::enable_if_t<std::is_same_v<T,
25 // std::string>>>
26 template <typename T, typename = std::enable_if_t<std::is_same<T,
27 std::string>::value>>
28 constexpr auto concat = sum<T>;
29
30 // Code from the assignment. The last parameter is needed
31 // because F is a type, even though it can be modified to only
32 // have 2 parameters, I believe this is easier on the head
33 template <typename T, typename F>
34 T reduce_sin(T* array, size_t n, F op) {
35     T result = array[0];
36     for (int i=1; i<n; ++i)
37         result = op(result, array[i]);
38     return result;
39 }
40
41 // I do not want to deal with promises or futures, so Im doing this.
42 /// Partially reduces the array from [start, end) and places it in
43 /// result.
44 /// @note The *result variable does not have a race condidtion as it
45 /// is only being written by one thread.
46 template <typename T, typename F>
47 void partial_reduce(T* start, T* end, T* result, F op) {
48     *result = reduce_sin(start, std::distance(start, end), op);
49 }
50
51 /// Parallel reduce
52 template <typename T, typename F>
53 T reduce_par (T* array, size_t n, size_t p, F op) {
54     // create nessicary structures
55     std::unique_ptr<T[]> data{ new T[p] };
56     std::vector<std::thread> threads{ };
57
58     // calculate parallelism
59     size_t depth = n / p;
60     T* start = array;
```

```

59     T* end = array + depth;
60
61     // the actual parallel code
62     for (size_t i = 0; i < p - 1; ++i) {
63         // The <T, F> is needed because partial_reduce is a templated function so
        we need
64         // to thread the <T, F> version of the function
65         threads.emplace_back(partial_reduce<T, F>, start, end, &data[i], op);
66         start = end;
67         end += depth;
68     }
69
70     // the last one (this thread will do it so we dont spin on waiting for them
    to join)
71     partial_reduce(start, array + n, &data[p - 1], op);
72
73     // wait for them to finish
74     for (auto& t : threads)
75         t.join();
76
77     // final sum
78     return reduce_sin(data.get(), p, op);
79 }
80
81
82 // options:
83 // 1: reduce<int, sum>
84 // 2: reduce<int, max>
85 // 3: reduce<std::string, concat>
86 // 4: reduce<float, sum>
87 // 5: reduce<float, max>
88 #define OPTION 1
89 #define MAX 500'000'007
90 // @warning, the only one that doesnt really work very well is option 4.
    Adding up
91 // more than 1mil values gives you two different answers and I'm mostly
    certain its
92 // because of FP arithmetic
93
94
95 #if OPTION == 1
96     #define NUMERIC
97     using NUM_TYPE = int;
98     template <typename T> constexpr auto op_t = sum<T>;
99 #elif OPTION == 2
100     #define NUMERIC
101     using NUM_TYPE = int;
102     template <typename T> constexpr auto op_t = max<T>;
103 #elif OPTION == 3
104     template <typename T> constexpr auto op_t = concat<T>;
105 #elif OPTION == 4
106     #define NUMERIC
107     using NUM_TYPE = float;
108     template <typename T> constexpr auto op_t = sum<T>;
109 #elif OPTION == 5
110     #define NUMERIC
111     using NUM_TYPE = float;
112     template <typename T> constexpr auto op_t = max<T>;
113 #else
114     #error "[E] Valid option not chosen"

```

```

115 #endif
116
117
118 int main() {
119
120 #ifdef NUMERIC
121
122     constexpr unsigned long long LG_NUM = MAX;
123     std::vector<NUM_TYPE> data{};
124     data.push_back(3); // this is a "valid" dummy data for max
125
126     for (unsigned long long i = 1; i < LG_NUM + 1; ++i) {
127         data.push_back(1); // sums of ones
128     }
129
130 #else
131
132     char temp[2] = { '\0', '\0' };
133     std::vector<std::string> data{};
134     for (int j = 0; j < 3; ++j) {
135         for (unsigned long long i = 0; i < 26; ++i) {
136             // 97-122
137             temp[0] = i + 97;
138             data.emplace_back(temp);
139         }
140     }
141
142 #endif
143
144     auto op = op_t<decltype(data)::value_type>;
145
146     auto begin1 = std::chrono::high_resolution_clock::now(); //// AHHHHHHH
147     auto accum1 = reduce_sin(data.data(), data.size(), op);
148     auto end1 = std::chrono::high_resolution_clock::now(); //// AHHHHHHH
149
150     std::chrono::duration<double, std::milli> elapse1{ end1 - begin1 };
151     std::cout << "Serial: " << elapse1.count() << std::endl;
152
153     auto begin2 = std::chrono::high_resolution_clock::now(); //// AHHHHHHH
154     auto accum2 = reduce_par(data.data(), data.size(), 8, op);
155     auto end2 = std::chrono::high_resolution_clock::now(); //// AHHHHHHH
156
157     std::chrono::duration<double, std::milli> elapse2{ end2 - begin2 };
158     std::cout << "Parallel: " << elapse2.count() << std::endl;
159
160     // Turns out that neither gnu g++ or clang++ supports the parallelism TS,
kms
161     // https://youtu.be/McJrc2uxbKI?t=571 MSVC ahead of the game *slow claps*
162     // nvm: https://godbolt.org/z/TCGaze The internet video lies
163     // auto begin3 = std::chrono::high_resolution_clock::now();
164     // double result = std::reduce(std::execution::par, sum.begin(),
sum.end());
165     // auto end3 = std::chrono::high_resolution_clock::now();
166
167     // std::chrono::duration<double, std::milli> elapse3{ end3 - begin3 };
168     // std::cout << "STL: " << elapse3.count();
169
170     // hehe https://gcc.gnu.org/onlinedocs/libstdc++/manual/parallel_mode.html
171     // after seeing the results, another lie. Why does the docs lie?????
172     // Serial: 1512.65

```

```
173 // Parallel: 351.756
174 // GNU: 1562.03
175 // The sum is: 500000010
176 // #if OPTION == 1
177 // auto begin4 = std::chrono::high_resolution_clock::now();
178 // double accum4 = __gnu_parallel::accumulate(data.begin(), data.end(), 0);
179 // auto end4 = std::chrono::high_resolution_clock::now();
180
181 // std::chrono::duration<double, std::milli> elapse4{ end4 - begin4 };
182 // std::cout << "GNU: " << elapse4.count() << std::endl;
183 // #endif
184
185 if (accum1 == accum2)
186     std::cout << "The result is: " << accum1 << std::endl;
187 else
188     std::cout << "The result were different, you screwed up..." << std::endl
<< accum1 << " " << accum2 << std::endl;
189 }
```